


1.Introduction



這次作業需要實做出一個雙層的 neural network，包括輸入層、兩層隱藏層與輸出層。其中由於訓練資料是二元分類問題，輸入層的維度為 2，輸出層的維度為 1。

當我們對神經元進行輸入後，經過內部迴歸模型對輸入的權重加乘，再經過 active function，便完成了該節點的輸出。該輸出會再傳給下一個神經元，作為該神經元的輸入值，如此一層層傳遞下去，直到最後一層的輸出層，產生預測結果，此過程稱為 Forward-Propagation。

neural network 會反覆由預測結果和真實結果之間的差距，對整個神經網路，由後面神經元至前層神經元的更新，此過程稱為 Back-Propagation。

透過這次作業的實做，能夠更清楚了解 Forward-Propagation 和 Back-Propagation 的工作原理。



2.Experiment setups

A. Sigmoid functions

在 neural network 中，若不使用 active function，neural network 即是以線性的方式組合運算，使用 active function，主要是利用非線性方程式，解決非線性問題。Sigmoid 函數是深度學習領域開始時使用頻率最高的 activation function，它是易於求導數之平滑函數，其函數和導函數圖形如下。並使用以下程式碼來實現。

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

$$\sigma'(x) = \frac{d}{dx} \left(\frac{1}{1 + e^{-x}} \right) = (-1) \cdot \frac{1}{(1 + e^{-x})^2} \cdot (-e^{-x}) = \frac{e^{-x}}{(1 + e^{-x})^2}$$

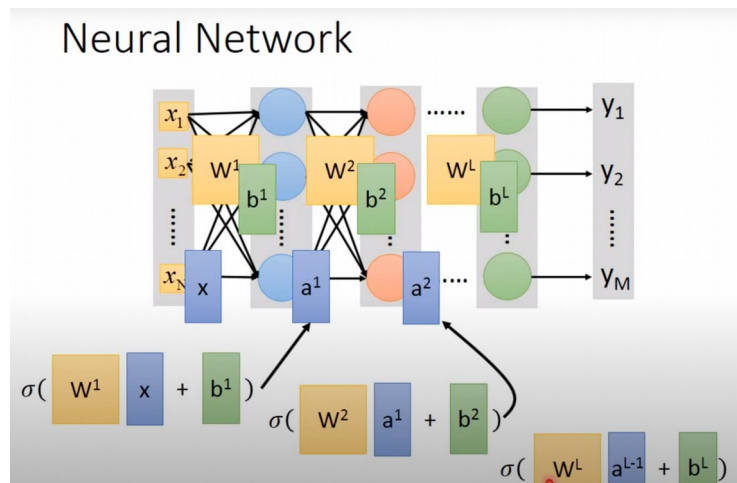


```
if self.activation_func == 'sigmoid':  
    return 1.0/(1.0+np.exp(-x))
```

```
if self.activation_func == 'sigmoid':  
    return np.multiply(x,1.0-x)
```

B. Neural network

假設上一層結點 i, j, k, \dots 等一些結點與本層的結點 w 有連接，那麼結點 w 的值就是通過上一層的 i, j, k 等結點以及對應的連接權值進行加權和運算，最後通過一個非線性函數（sigmoid 等函數），最後得到的結果就是本層結點 w 的輸出。最終不斷的通過這種方法一層層的運算，得到輸出層結果。



程式實做（我沒有寫 bias 項）

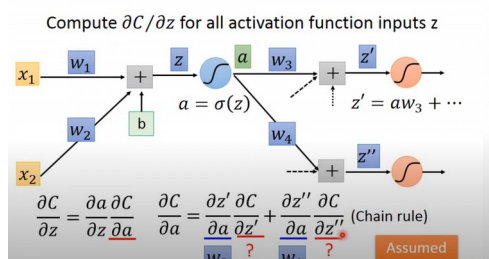
```
def forward(self, x, W):
    Z = list([])
    A = list([])
    for l in range(self.layers+1):
        if not l==0: x = A[-1]
        Z.append( np.dot(x, W[l]) )
        A.append(np.array([self.activation(item) for item in Z[-1] ], dtype=np.float128))

    pred_y = A[-1]
    return Z, A, pred_y
```


C. Backpropagation

Backpropagation 是誤差反向傳播的簡稱，是一種與最佳化方法（如梯度下降法）結合使用的方法，透過微積分的連鎖律，我們可以計算出損失函數對每一個節點的梯度，再進一步算出對每個權重值的梯度。


Backpropagation – Backward pass



Backpropagation – Summary



Backpropagation – Backward pass



程式實做

```
def back(self, W, A, pred_y, x, y, num_of_nodes):
    dW = init_parameters_zeros(num_of_nodes)
    dZ = list([])

    for k in range(self.layers, -1, -1):
        if k == self.layers:
            tmp_dZlist = []
            for i in range(len(y)):
                tmp_dZlist.append((y[i] - pred_y[i]) * self.derivative_activation(A[self.layers][i]))
            dZ.insert(0, np.array(tmp_dZlist, dtype=np.float128))
        else:
            tmp_dZlist = list([])
            for i in range(len(A[k])):
                dZtmp = 0
                for j in range(len(dZ[0])):
                    dZtmp += dZ[0][j] * W[k+1][i][j]
                dZtmp = dZtmp * self.derivative_activation(A[k][i])
                tmp_dZlist.append(dZtmp)
            dZ.insert(0, np.array(tmp_dZlist, dtype=np.float128))

    for i in range(self.layers+1):
        for j in range(len(W[i])):
            for k in range(len(W[i][j])):
                if not i==0:
                    dW[i][j][k] = dZ[i-1][j] * dZ[i][k]
                else:
                    dW[i][j][k] = x[j] * dZ[i][k]



    return dW
```

3.Results of your testing

(1) linear

程式執行方法：

```
python3 hw1.py --task linear --batch_size 100 --lr 0.5 --epochs 100000 --optimizers  
GD --N 100 --hidden_units 2 --activation_func sigmoid
```





```
20:45:22 epoch= 85000 Loss= 0.20759575363409612489  
20:45:26 epoch= 85500 Loss= 0.20751433578804313922  
20:45:31 epoch= 86000 Loss= 0.2074324098918373797  
20:45:35 epoch= 86500 Loss= 0.20734994077144711426  
20:45:40 epoch= 87000 Loss= 0.20726689254108546592  
20:45:45 epoch= 87500 Loss= 0.20718322857602706761  
20:45:50 epoch= 88000 Loss= 0.20709891148381432316  
20:45:55 epoch= 88500 Loss= 0.20701390307351310464  
20:46:00 epoch= 89000 Loss= 0.20692816432270213292  
20:46:05 epoch= 89500 Loss= 0.20684165534194422024  
20:46:10 epoch= 90000 Loss= 0.20675433533660602527  
20:46:14 epoch= 90500 Loss= 0.20666616256608448429  
20:46:19 epoch= 91000 Loss= 0.20657709430078496037  
20:46:23 epoch= 91500 Loss= 0.20648708677760472612  
20:46:28 epoch= 92000 Loss= 0.2063960951552359438  
20:46:33 epoch= 92500 Loss= 0.20630407347134864885  
20:46:38 epoch= 93000 Loss= 0.2062109746046827275  
20:46:43 epoch= 93500 Loss= 0.20611675024630559485  
20:46:48 epoch= 94000 Loss= 0.20602135088581424022  
20:46:53 epoch= 94500 Loss= 0.20592472582010517174  
20:46:57 epoch= 95000 Loss= 0.20582682319451994184  
20:47:02 epoch= 95500 Loss= 0.20572759008869347733  
20:47:07 epoch= 96000 Loss= 0.20562697266225289524  
20:47:11 epoch= 96500 Loss= 0.20552491637855832733  
20:47:16 epoch= 97000 Loss= 0.20542136632780940217  
20:47:21 epoch= 97500 Loss= 0.20531626767385346371  
20:47:26 epoch= 98000 Loss= 0.20520956625162885572  
20:47:30 epoch= 98500 Loss= 0.20510120934396494103  
20:47:35 epoch= 99000 Loss= 0.20499114666694482534
```

```
0.49974551572716888035  
0.49974497156744092995  
0.10576699123627181962  
0.001810226164572917604  
6.668015379984051555e-05  
0.49970775629922699687  
0.00046363212100520100142  
7.417679848051290241e-08  
6.630279650598467335e-08  
0.4992948417143853214  
0.08890652394399669987  
6.649099278488791028e-11  
0.48076408631514147793  
0.4996107690506225482  
0.4997448624576181588  
0.0023031342467738315717  
0.49588122929799338326  
2.5141962195083000392e-07  
0.49969900970342090107  
1.1987831969783921778e-08  
0.49969643066571512097  
0.49971527392058805403  
0.21369239233349767478  
0.0002791390605866182078  
0.4994672244723581556  
0.08910685118595268023  
7.363775255069901435e-10  
0.4997461894694492084  
accuracy : 1.0
```

(2) xor

```
python3 hw1.py --task xor --batch_size 21 --lr 0.02 --epochs 10000 --optimizers Adam --hidden_units 30 --activation_func sigmoid
```




```
02:36:31 epoch= 0 Loss= 0.49549586737734632483
02:36:56 epoch= 500 Loss= 0.35675747615130063884
02:37:19 epoch= 1000 Loss= 0.31335750371031477102
02:37:42 epoch= 1500 Loss= 0.2913509120607740668
02:38:05 epoch= 2000 Loss= 0.28571249537389593445
02:38:29 epoch= 2500 Loss= 0.28539397559191799443
02:38:52 epoch= 3000 Loss= 0.28388759153271947333
02:39:16 epoch= 3500 Loss= 0.27708786579206244572
02:39:40 epoch= 4000 Loss= 0.2648575360102687917
02:40:05 epoch= 4500 Loss= 0.26193728070663100513
02:40:31 epoch= 5000 Loss= 0.2619132865129534643
02:40:59 epoch= 5500 Loss= 0.2619105366670904583
02:41:26 epoch= 6000 Loss= 0.26190936036489775204
02:41:50 epoch= 6500 Loss= 0.26190867740328521
02:42:14 epoch= 7000 Loss= 0.2619082203111979733
02:42:38 epoch= 7500 Loss= 0.2619078879460048737
02:43:02 epoch= 8000 Loss= 0.26190763274850152324
02:43:27 epoch= 8500 Loss= 0.26190742910257083003
02:43:51 epoch= 9000 Loss= 0.26190726185843124797
02:44:15 epoch= 9500 Loss= 0.26190712142016237338
```

```
0.5
0.5
4.703092029651021681e-05
0.5
6.2920828600767705106e-454
0.5
8.8821978337533592615e-512
0.5
3.3659700338326600073e-514
0.49999999861078816426
2.6693911799201342352e-514
2.6437772360392170901e-514
0.4999999986132642944
2.6427174629028789133e-514
0.5
2.6426734031988354367e-514
0.5
2.6426715710629240175e-514
0.5
2.64267149487643784e-514
0.5
accuracy : 0.9523809523809523
```

4. Discussion



A. Try different learning rates

將 learning rates 調大，可以加快一開始 Loss 的收斂速度。但有時候 learning rates 過大，在某些地方會卡住，使 Loss 在兩個值間震盪下不去。



B. Try different numbers of hidden units

因為非線性的 xor 較難訓練。當我用一層只有 2 個 hidden units 時，雖然能夠將 Loss 降低到很小，但是它僅將預測正確的逼近 0 和 1，預測結果仍然偏線性，約 1/4 預測錯誤。使用多一些 hidden units 才能有較好預測結果。



```
05:58:00 epoch= 286500 Loss= 0.00054792735293966266955
05:58:01 epoch= 287000 Loss= 0.0005476785417664186075
05:58:02 epoch= 287500 Loss= 0.00054743028168533928043
05:58:02 epoch= 288000 Loss= 0.0005471825705130574826
05:58:03 epoch= 288500 Loss= 0.00054693540607944157407
05:58:04 epoch= 289000 Loss= 0.0005466887862273539076
05:58:05 epoch= 289500 Loss= 0.00054644270880967181567
05:58:05 epoch= 290000 Loss= 0.0005461971716936820091
05:58:06 epoch= 290500 Loss= 0.00054595217275700273315
05:58:07 epoch= 291000 Loss= 0.0005457077098917001413
05:58:08 epoch= 291500 Loss= 0.00054546378099891092985
05:58:09 epoch= 292000 Loss= 0.0005452203839933364805
05:58:09 epoch= 292500 Loss= 0.0005449775168006355246
05:58:10 epoch= 293000 Loss= 0.00054473517735849021666
05:58:11 epoch= 293500 Loss= 0.00054449336361611060425
05:58:12 epoch= 294000 Loss= 0.00054425207353421204307
05:58:12 epoch= 294500 Loss= 0.00054401130508464122736
05:58:13 epoch= 295000 Loss= 0.0005437710562505707861
05:58:14 epoch= 295500 Loss= 0.0005435313250263353106
05:58:15 epoch= 296000 Loss= 0.0005432921094177109404
05:58:15 epoch= 296500 Loss= 0.00054305340744099277674
05:58:16 epoch= 297000 Loss= 0.0005428152171247081792
05:58:17 epoch= 297500 Loss= 0.00054257753650647374207
05:58:18 epoch= 298000 Loss= 0.00054234036363682570164
05:58:19 epoch= 298500 Loss= 0.00054210369657478776193
05:58:19 epoch= 299000 Loss= 0.0005418675333919530598
05:58:20 epoch= 299500 Loss= 0.00054163187216857342824
```

```
0.0059553689890424601597
1.0
4.3658388576987397207e-93
1.0
4.3658388576987397207e-93
1.0
4.3658388576987397207e-93
1.0
4.3658388576987397207e-93
1.0
4.3658388576987397207e-93
1.0
4.3658388576987397207e-93
4.365837015258929447e-93
4.365837015258929447e-93
4.365837015258929447e-93
4.365837015258929447e-93
4.365837015258929447e-93
4.365837015258929447e-93
4.365837015258929447e-93
4.365837015258929447e-93
4.365837015258929447e-93
4.365837015258929447e-93
4.365837015258929447e-93
accuracy : 0.7619047619047619
```

C. Try without activation functions

計算過程中直接炸裂到無限大了。

5. Extra

A .Implement different optimizers

```
if self.optimizers == 'GD':
    new_W = init_parameters_zeros(self.num_of_nodes)
    for i in range(self.layers+1):
        for j in range(len(W[i])):
            for k in range(len(W[i][j])):
                new_W[i][j][k] = W[i][j][k] + self.lr * dW[i][j][k]
    return new_W
```

```
if self.optimizers == 'Momentum':
    global vt_last
    vt = init_parameters_zeros(self.num_of_nodes)
    new_W = init_parameters_zeros(self.num_of_nodes)
    beta = 0.9
    try : _ = vt_last
    except NameError: vt_last = init_parameters_zeros(self.num_of_nodes)

    for i in range(self.layers+1):
        for j in range(len(W[i])):
            for k in range(len(W[i][j])):
                vt[i][j][k] = beta * vt_last[i][j][k] + self.lr * dW[i][j][k]
                new_W[i][j][k] = W[i][j][k] + vt[i][j][k]
                vt_last[i][j][k] = vt[i][j][k]

    return new_W
```

```
if self.optimizers == 'AdaGrad':
    epsilon = 0.00000001
    n = 0
    for i in range(self.layers+1):
        for j in range(len(W[i])):
            for k in range(len(W[i][j])):
                n += dW[i][j][k] * dW[i][j][k]

    new_W = init_parameters_zeros(self.num_of_nodes)
    for i in range(self.layers+1):
        for j in range(len(W[i])):
            for k in range(len(W[i][j])):
                new_W[i][j][k] = W[i][j][k] + self.lr * dW[i][j][k] * (n+epsilon)**(-0.5)
    return new_W
```

```

if self.optimizers == 'Adam':
    global vt_old
    global mt_old
    vt = init_parameters_zeros(self.num_of_nodes)
    mt = init_parameters_zeros(self.num_of_nodes)
    vt_hat = init_parameters_zeros(self.num_of_nodes)
    mt_hat = init_parameters_zeros(self.num_of_nodes)
    new_W = init_parameters_zeros(self.num_of_nodes)
    beta = 0.9
    epsilon = 0.00000001

    try : __, __ = vt_old, mt_old
    except NameError:
        vt_old = init_parameters_zeros(self.num_of_nodes)
        mt_old = init_parameters_zeros(self.num_of_nodes)
    for i in range(self.layers+1):
        for j in range(len(W[i])):
            for k in range(len(W[i][j])):
                vt[i][j][k] = beta * vt_old[i][j][k] + (1.0 - beta) * dW[i][j][k] * dW[i][j][k]
                mt[i][j][k] = beta * mt_old[i][j][k] + (1.0 - beta) * dW[i][j][k]

                vt_hat[i][j][k] = vt[i][j][k]/(1.0 - beta)
                mt_hat[i][j][k] = mt[i][j][k]/(1.0 - beta)

                new_W[i][j][k] = W[i][j][k] + self.lr * (mt_hat[i][j][k] / ((vt_hat[i][j][k])**0.5+epsilon))
    return new_W

```

B. Implement different activation functions

```

def activation(self,x):
    if self.activation_func == 'sigmoid':
        return 1.0/(1.0+np.exp(-x))
    if self.activation_func == 'None':
        return 1.0 * x
    if self.activation_func == 'tanh':
        return np.tanh(x)
    if self.activation_func == "ReLU":
        return np.maximum(0, x)

def derivative_activation(self,x):
    if self.activation_func == 'sigmoid':
        return np.multiply(x,1.0-x)
    if self.activation_func == 'None':
        return 1.0
    if self.activation_func == 'tanh':
        return 1.0 - x ** 2
    if self.activation_func == "ReLU":
        return 1.0 * (x > 0)

```